

car model

May 9, 2018

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from skimage.io import imread
from skimage.transform import downscale_local_mean
from os.path import join
from tqdm import tqdm_notebook
import cv2
from sklearn.model_selection import train_test_split

In [5]: input_folder = join('input')

df_mask = pd.read_csv(join(input_folder, 'train_masks.csv'), usecols=['img'])
ids_train = df_mask['img'].map(lambda s: s.split('_')[0]).unique()

imgs_idx = list(range(1, 17))

In [6]: load_img = lambda im, idx: imread(join(input_folder, 'traincar', '{}_{:02d}.jpg'.format(
load_mask = lambda im, idx: imread(join(input_folder, 'train_masks', '{}_{:02d}_mask.gif'.format(
resize = lambda im: downscale_local_mean(im, (4,4) if im.ndim==2 else (4,4,1))
mask_image = lambda im, mask: (im * np.expand_dims(mask, 2))

In [7]: num_train = 32 # len(ids_train)

# Load data for position id=1
X = np.empty((num_train, 320, 480, 12), dtype=np.float32)
y = np.empty((num_train, 320, 480, 1), dtype=np.float32)

with tqdm_notebook(total=num_train) as bar:
    idx = 1 # Rotation index
    for i, img_id in enumerate(ids_train[:num_train]):
        imgs_id = [resize(load_img(img_id, j)) for j in imgs_idx]
        # Input is image + mean image per channel + std image per channel
        X[i, ..., :9] = np.concatenate([imgs_id[idx-1], np.mean(imgs_id, axis=0), np.std(
        y[i] = resize(np.expand_dims(load_mask(img_id, idx), 2)) / 255.
        del imgs_id # Free memory
        bar.update()
```

```
HBox(children=(IntProgress(value=0, max=32), HTML(value='')))
```

```
In [8]: X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
In [9]: y_train_mean = y_train.mean(axis=0)
        y_train_std = y_train.std(axis=0)
        y_train_min = y_train.min(axis=0)
```

```
y_features = np.concatenate([y_train_mean, y_train_std, y_train_min], axis=2)
```

```
X_train[:, ..., -3:] = y_features
X_val[:, ..., -3:] = y_features
```

```
In [10]: X_mean = X_train.mean(axis=(0,1,2), keepdims=True)
        X_std = X_train.std(axis=(0,1,2), keepdims=True)
```

```
X_train -= X_mean
X_train /= X_std
```

```
X_val -= X_mean
X_val /= X_std
```

```
In [11]: from keras.layers import Conv2D
        from keras.models import Sequential
        import keras.backend as K
```

```
model = Sequential()
model.add( Conv2D(16, 3, activation='relu', padding='same', input_shape=(320, 480, 12))
model.add( Conv2D(32, 3, activation='relu', padding='same') )
model.add( Conv2D(1, 5, activation='sigmoid', padding='same') )
```

```
/home/nikhil/anaconda3/lib/python3.6/site-packages/h5py/__init__.py:36: FutureWarning: Conversion
from ._conv import register_converters as _register_converters
Using TensorFlow backend.
```

```
In [12]: from keras.optimizers import Adam
        from keras.losses import binary_crossentropy
```

```
smooth = 1.
```

```
# From here: https://github.com/jocicmarko/ultrasound-nerve-segmentation/blob/master/tr
def dice_coef(y_true, y_pred):
    y_true_f = K.flatten(y_true)
    y_pred_f = K.flatten(y_pred)
```

```

        intersection = K.sum(y_true_f * y_pred_f)
        return (2. * intersection + smooth) / (K.sum(y_true_f) + K.sum(y_pred_f) + smooth)

def bce_dice_loss(y_true, y_pred):
    return 0.5 * binary_crossentropy(y_true, y_pred) - dice_coef(y_true, y_pred)

model.compile(Adam(lr=1e-3), bce_dice_loss, metrics=['accuracy', dice_coef])

In [14]: history = model.fit(X_train, y_train, epochs=15, validation_data=(X_val, y_val), batch_

Train on 25 samples, validate on 7 samples
Epoch 1/15
- 18s - loss: -5.1759e-02 - acc: 0.7565 - dice_coef: 0.3500 - val_loss: -2.2646e-01 - val_acc:
Epoch 2/15
- 18s - loss: -3.4551e-01 - acc: 0.8516 - dice_coef: 0.5267 - val_loss: -4.6521e-01 - val_acc:
Epoch 3/15
- 18s - loss: -5.7210e-01 - acc: 0.9391 - dice_coef: 0.6682 - val_loss: -6.6045e-01 - val_acc:
Epoch 4/15
- 19s - loss: -7.2053e-01 - acc: 0.9472 - dice_coef: 0.7882 - val_loss: -7.3929e-01 - val_acc:
Epoch 5/15
- 18s - loss: -7.6174e-01 - acc: 0.9493 - dice_coef: 0.8345 - val_loss: -7.7313e-01 - val_acc:
Epoch 6/15
- 18s - loss: -7.8110e-01 - acc: 0.9517 - dice_coef: 0.8572 - val_loss: -7.7548e-01 - val_acc:
Epoch 7/15
- 19s - loss: -8.0069e-01 - acc: 0.9572 - dice_coef: 0.8651 - val_loss: -8.0232e-01 - val_acc:
Epoch 8/15
- 18s - loss: -8.1617e-01 - acc: 0.9609 - dice_coef: 0.8717 - val_loss: -8.0282e-01 - val_acc:
Epoch 9/15
- 18s - loss: -8.1364e-01 - acc: 0.9604 - dice_coef: 0.8746 - val_loss: -8.0195e-01 - val_acc:
Epoch 10/15
- 18s - loss: -8.2044e-01 - acc: 0.9618 - dice_coef: 0.8714 - val_loss: -8.0899e-01 - val_acc:
Epoch 11/15
- 18s - loss: -8.2184e-01 - acc: 0.9617 - dice_coef: 0.8778 - val_loss: -8.0989e-01 - val_acc:
Epoch 12/15
- 18s - loss: -8.2939e-01 - acc: 0.9635 - dice_coef: 0.8796 - val_loss: -8.1681e-01 - val_acc:
Epoch 13/15
- 18s - loss: -8.3228e-01 - acc: 0.9636 - dice_coef: 0.8845 - val_loss: -8.1460e-01 - val_acc:
Epoch 14/15
- 18s - loss: -8.2932e-01 - acc: 0.9628 - dice_coef: 0.8851 - val_loss: -8.2089e-01 - val_acc:
Epoch 15/15
- 18s - loss: -8.3372e-01 - acc: 0.9645 - dice_coef: 0.8818 - val_loss: -8.2552e-01 - val_acc:

In [15]: plt.imshow(y_pred > 0.5, cmap='gray')

```

NameError


Traceback (most recent call last)

```
<ipython-input-15-faab2febc72c> in <module>()
----> 1 plt.imshow(y_pred > 0.5, cmap='gray')
```

NameError: name 'y_pred' is not defined

```
In [16]: pd.DataFrame(history.history)[['dice_coef', 'val_dice_coef']].plot()
```

```
Out[16]: <matplotlib.axes._subplots.AxesSubplot at 0x7f83c52c4cf8>
```



```
In [17]: idx = 0
         x = X_val[idx]

         fig, ax = plt.subplots(5,3, figsize=(16, 16))
         ax = ax.ravel()

         cmaps = ['Reds', 'Greens', 'Blues']
         for i in range(x.shape[-1]):
             ax[i].imshow(x[...,i], cmap='gray') #cmaps[i%3])
             ax[i].set_title('channel {}'.format(i))

         ax[-3].imshow((x[..., :3] * X_std[0, ..., :3] + X_mean[0, ..., :3]) / 255.)
         ax[-3].set_title('X')
```

```
y_pred = model.predict(x[None]).squeeze()
ax[-1].imshow(y_pred, cmap='gray')
ax[-1].set_title('y_pred')
```

```
Out[18]: <matplotlib.image.AxesImage at 0x7f83b277d668>
```

